Name: Kaushik Kotian Roll no.:29 Div:D20B Batch:B

AI and DS-II ASSIGNMENT - 2

1. Explain Restricted Boltzmann Machine(RBM).

A Restricted Boltzmann Machine (RBM) is a generative stochastic neural network that can learn a probability distribution over its set of inputs. RBMs are used as building blocks for deep learning models such as deep belief networks and deep autoencoders. They consist of two layers: a visible layer (input layer) and a hidden layer, with no connections between the neurons in the same layer (hence "restricted").

RBMs are used primarily for dimensionality reduction, feature learning, and generative modeling. They learn a joint probability distribution of the input data by adjusting the weights between the visible and hidden layers, through an unsupervised learning process. RBMs utilize an energy-based model, where the objective is to minimize the energy function, which measures how well a set of inputs fits the learned distribution. The hidden units in RBMs help to capture the underlying patterns in the data by modeling correlations between visible units.

The learning in RBMs typically occurs using contrastive divergence, a simplified and fast approximation of the more computationally expensive Gibbs sampling technique. Once trained, RBMs can generate new data points similar to those from the training data, making them useful for tasks like collaborative filtering and generative modeling.

2. Explain comparative analysis of different ML techniques.

Machine Learning (ML) encompasses various techniques, each having strengths and weaknesses depending on the nature of the problem. Below is a comparative analysis of some of the most common techniques:

- **Linear Regression vs. Decision Trees**: Linear regression is a parametric algorithm that assumes a linear relationship between input features and output, which limits its flexibility. Decision trees, on the other hand, are non-parametric and can capture non-linear relationships. However, decision trees are prone to overfitting unless pruned, whereas linear regression tends to be more stable but less expressive.
- **Support Vector Machines (SVM) vs. Neural Networks**: SVMs are effective for small- to medium-sized datasets with well-defined margins

between classes. They work well for binary classification and can handle high-dimensional spaces efficiently. Neural networks, particularly deep learning models, excel in large datasets with complex, non-linear relationships but require more computational power and tuning (such as optimizing learning rates and network architectures).

- **Random Forest vs. Gradient Boosting**: Both are ensemble methods. Random forests create multiple decision trees in parallel and combine their results, making them robust to overfitting and efficient in training. Gradient boosting builds trees sequentially, improving on previous trees' mistakes, which generally makes it more accurate but also slower and prone to overfitting if not regulated.
- **K-Means Clustering vs. DBSCAN**: K-means clustering is easy to implement but requires specifying the number of clusters ahead of time and struggles with non-spherical clusters. DBSCAN, a density-based clustering method, automatically detects the number of clusters and can handle arbitrary shapes, making it more suitable for complex data structures.

Overall, the choice of algorithm depends on dataset size, problem complexity, and the need for interpretability vs. flexibility.

3. Explain Large Scale Visual Recognition Challenge.

The Large Scale Visual Recognition Challenge (LSVRC) is a competition in the field of computer vision, which aims to evaluate algorithms for image classification, object detection, and other visual recognition tasks. It is part of the broader ImageNet project and has helped drive significant advances in the development of deep learning models.

The challenge involves classifying and detecting objects in images at a large scale, using the ImageNet dataset, which contains millions of labeled images. LSVRC offers multiple tasks each year, such as image classification (assigning a label to an image from a set of categories), object detection (detecting the location of objects in an image), and object localization (identifying the bounding boxes of objects).

The competition has driven substantial innovations in computer vision, with notable achievements including the emergence of convolutional neural networks (CNNs), especially since the success of AlexNet in 2012. These developments have had profound implications for industries like autonomous driving, healthcare, and image search technologies.

4. Explain Image Net Large Scale Visual Recognition Challenge (ILSVRC).

The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) is a specific instance of the LSVRC, based on the ImageNet dataset, which contains over 14 million images classified into 20,000 categories. Since its inception in 2010, ILSVRC has become a benchmark for evaluating the performance of object recognition algorithms.

Participants are required to develop models capable of classifying images into 1,000 object categories. The challenge offers multiple tasks:

- **Image classification**: Assign a single label to each image.
- **Object localization**: Identify the object within an image with bounding boxes.
- **Object detection**: Detect multiple objects in an image with bounding boxes and assign class labels.

AlexNet's groundbreaking win in the 2012 ILSVRC competition popularized CNNs, and subsequent models like VGGNet, GoogLeNet, ResNet, and EfficientNet have set new performance standards, often surpassing human accuracy in image classification tasks. ILSVRC has catalyzed advancements in transfer learning, network architectures, and optimization techniques.

5. Explain learning rate in neural network model.

In neural networks, the learning rate is a hyperparameter that controls how much the model's weights are adjusted during training with respect to the loss gradient. It is a crucial factor that affects the convergence of the model.

- **Too high learning rate**: If the learning rate is too large, the model may fail to converge because it overshoots the optimal weights during training. This leads to poor accuracy as the model oscillates around the minimum of the loss function without settling.
- **Too low learning rate**: On the other hand, if the learning rate is too small, the training process becomes extremely slow, and the model might get stuck in local minima, failing to reach the global minimum.

A well-chosen learning rate balances the speed and accuracy of convergence. In practice, learning rates are often adjusted dynamically during training using techniques like **learning rate decay**, **momentum**, or **adaptive methods** like **Adam**, which automatically adjust the learning rate for each parameter.

6. Draw architecture of auto encoders.